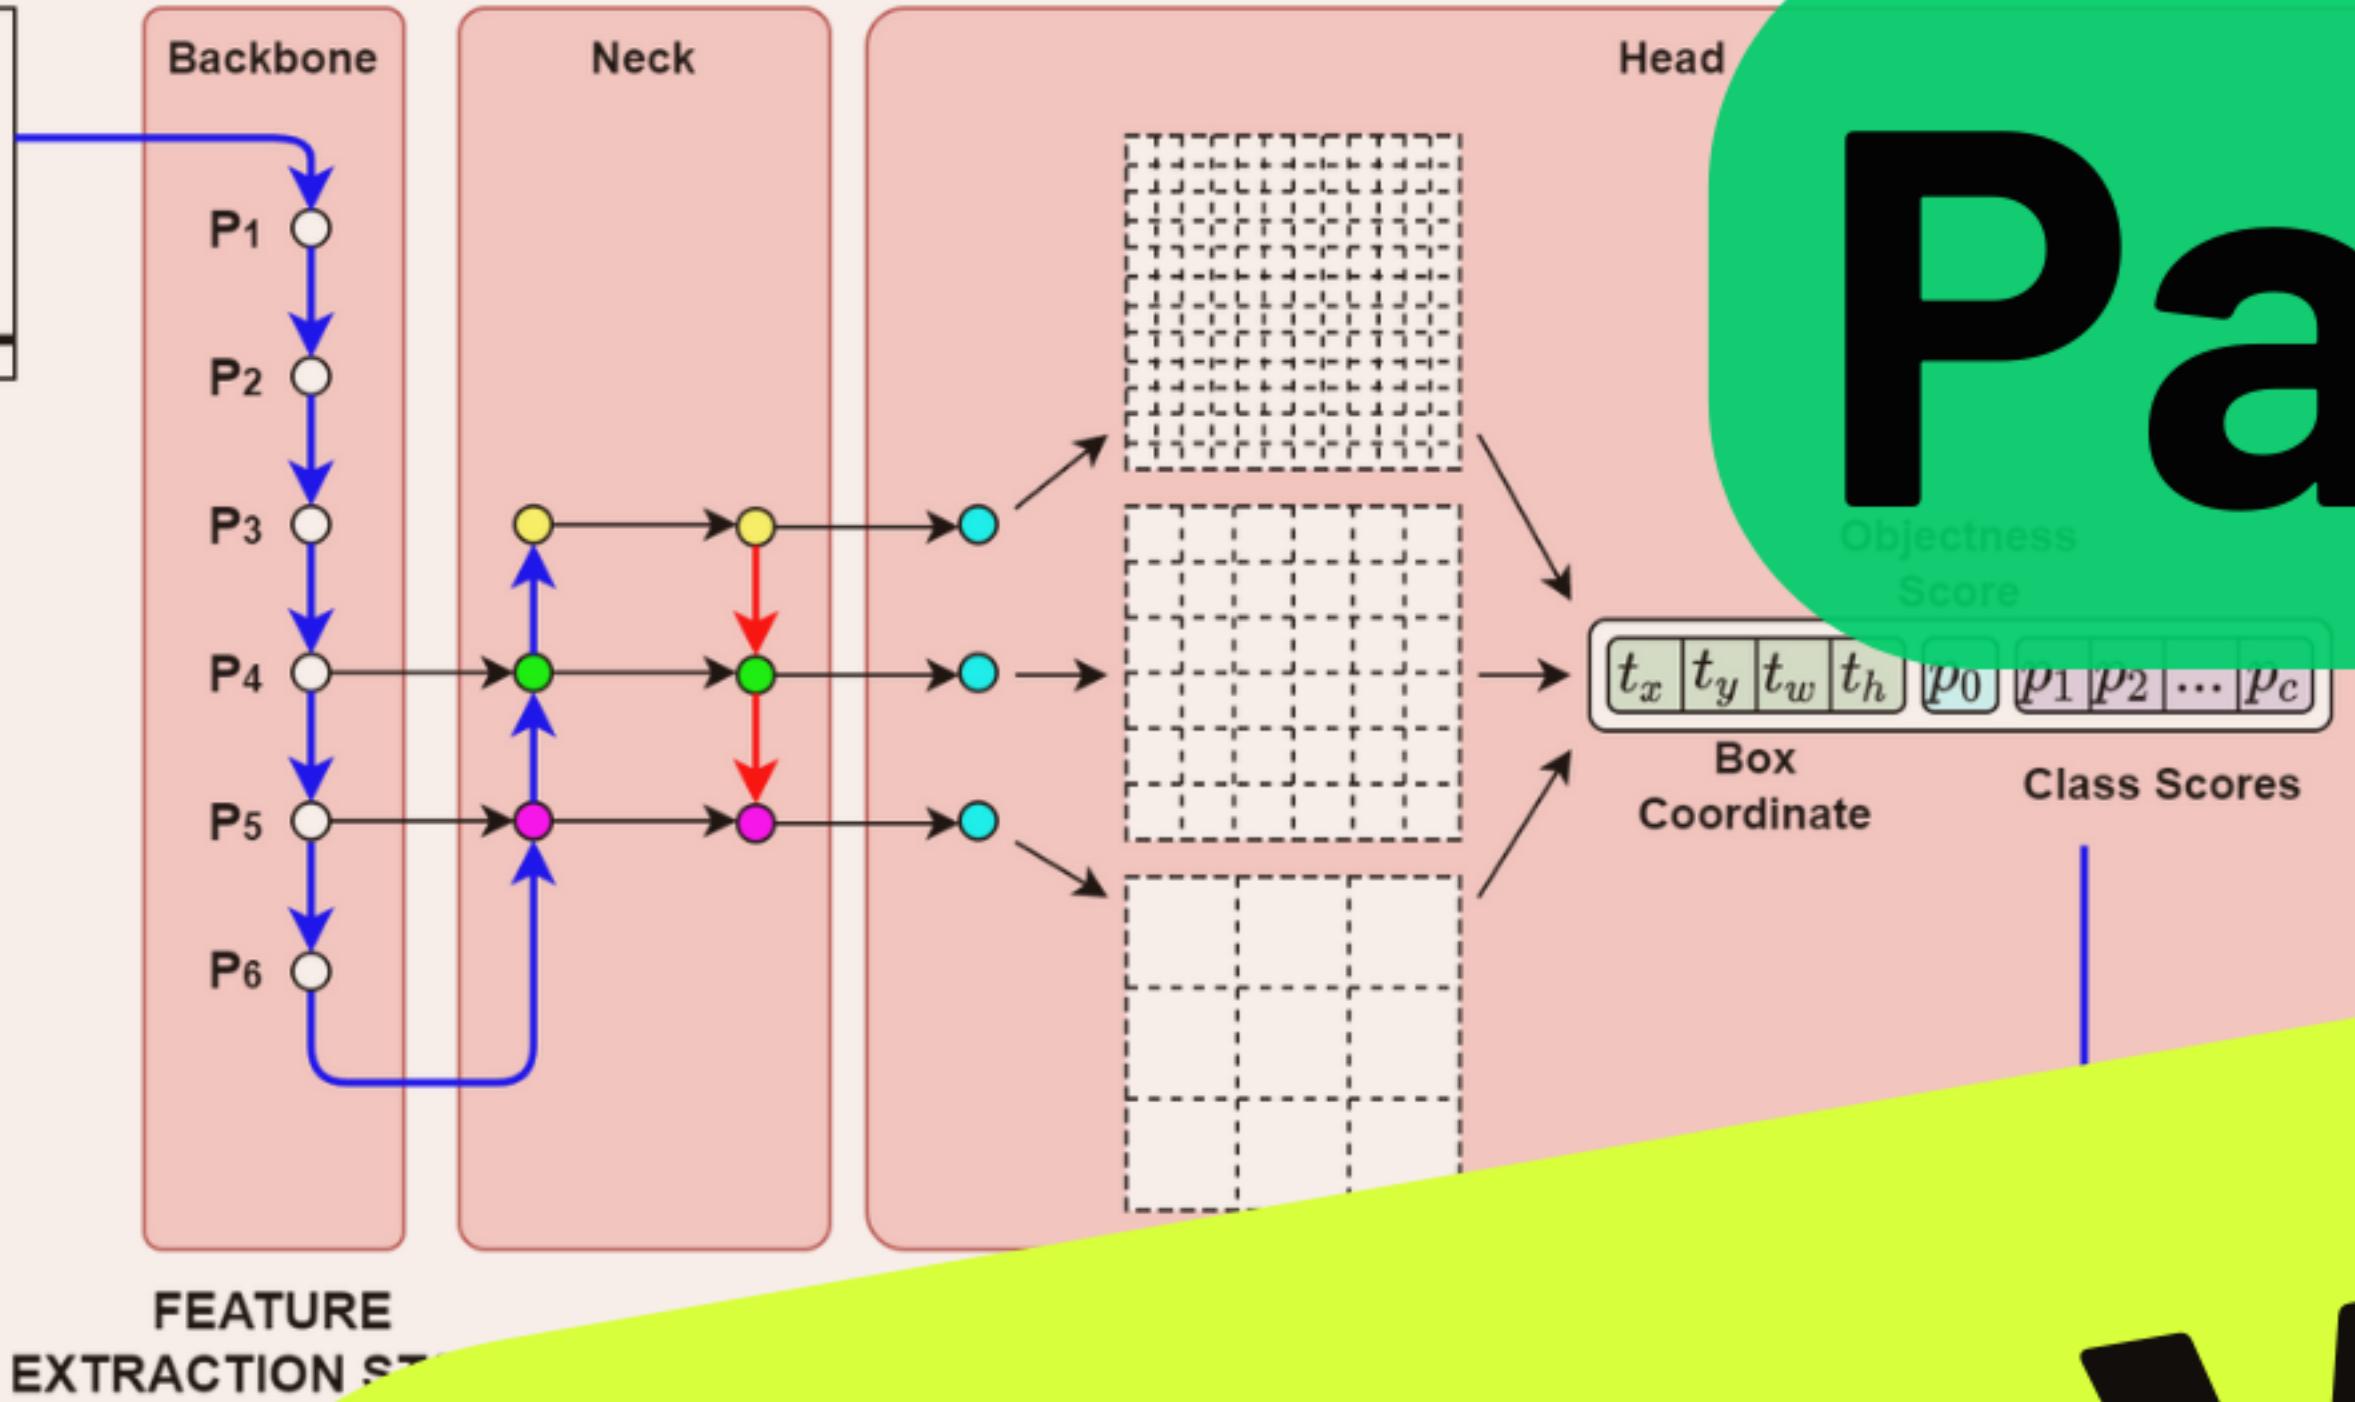
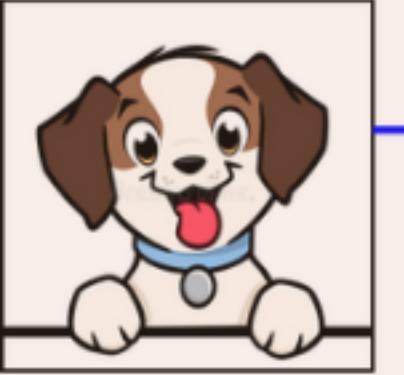
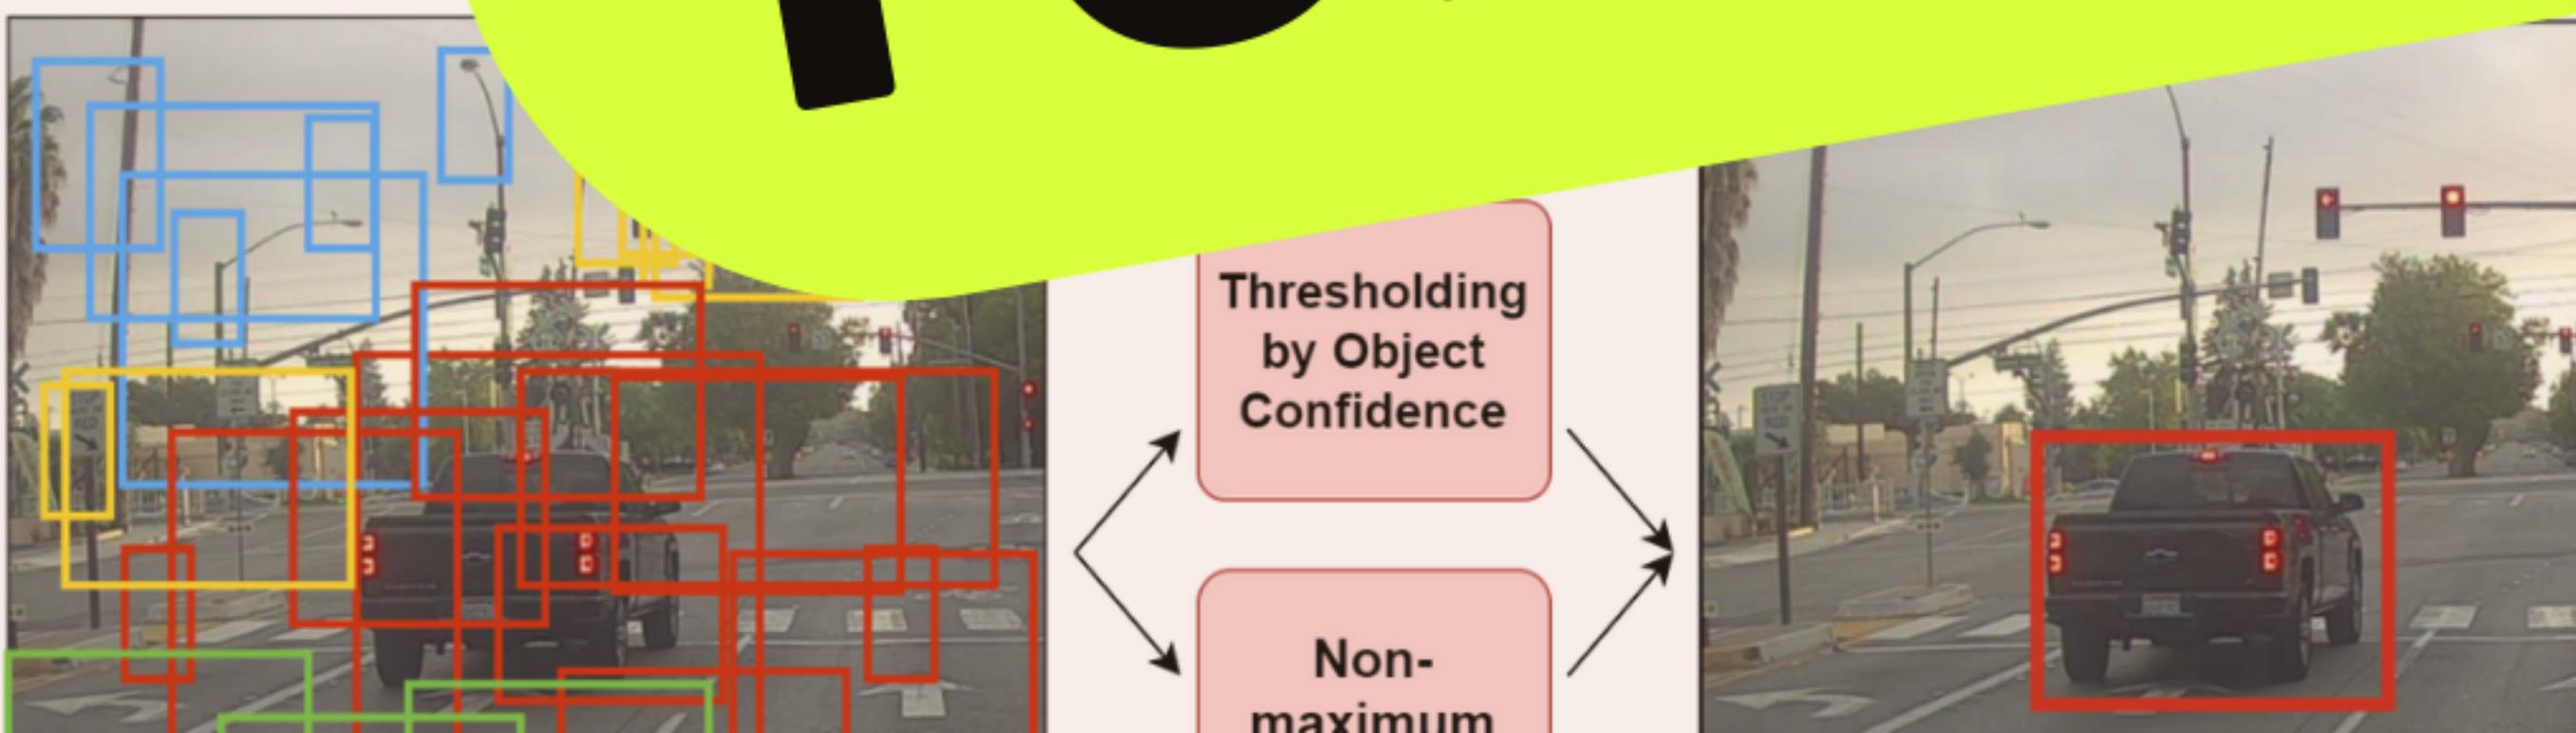


RICH SEMANTIC
EXTRACTION METHODS



YOLO-v4



Data Augmentation

- Random erase
- CutOut
- MixUp
- CutMix
- Style transfer GAN

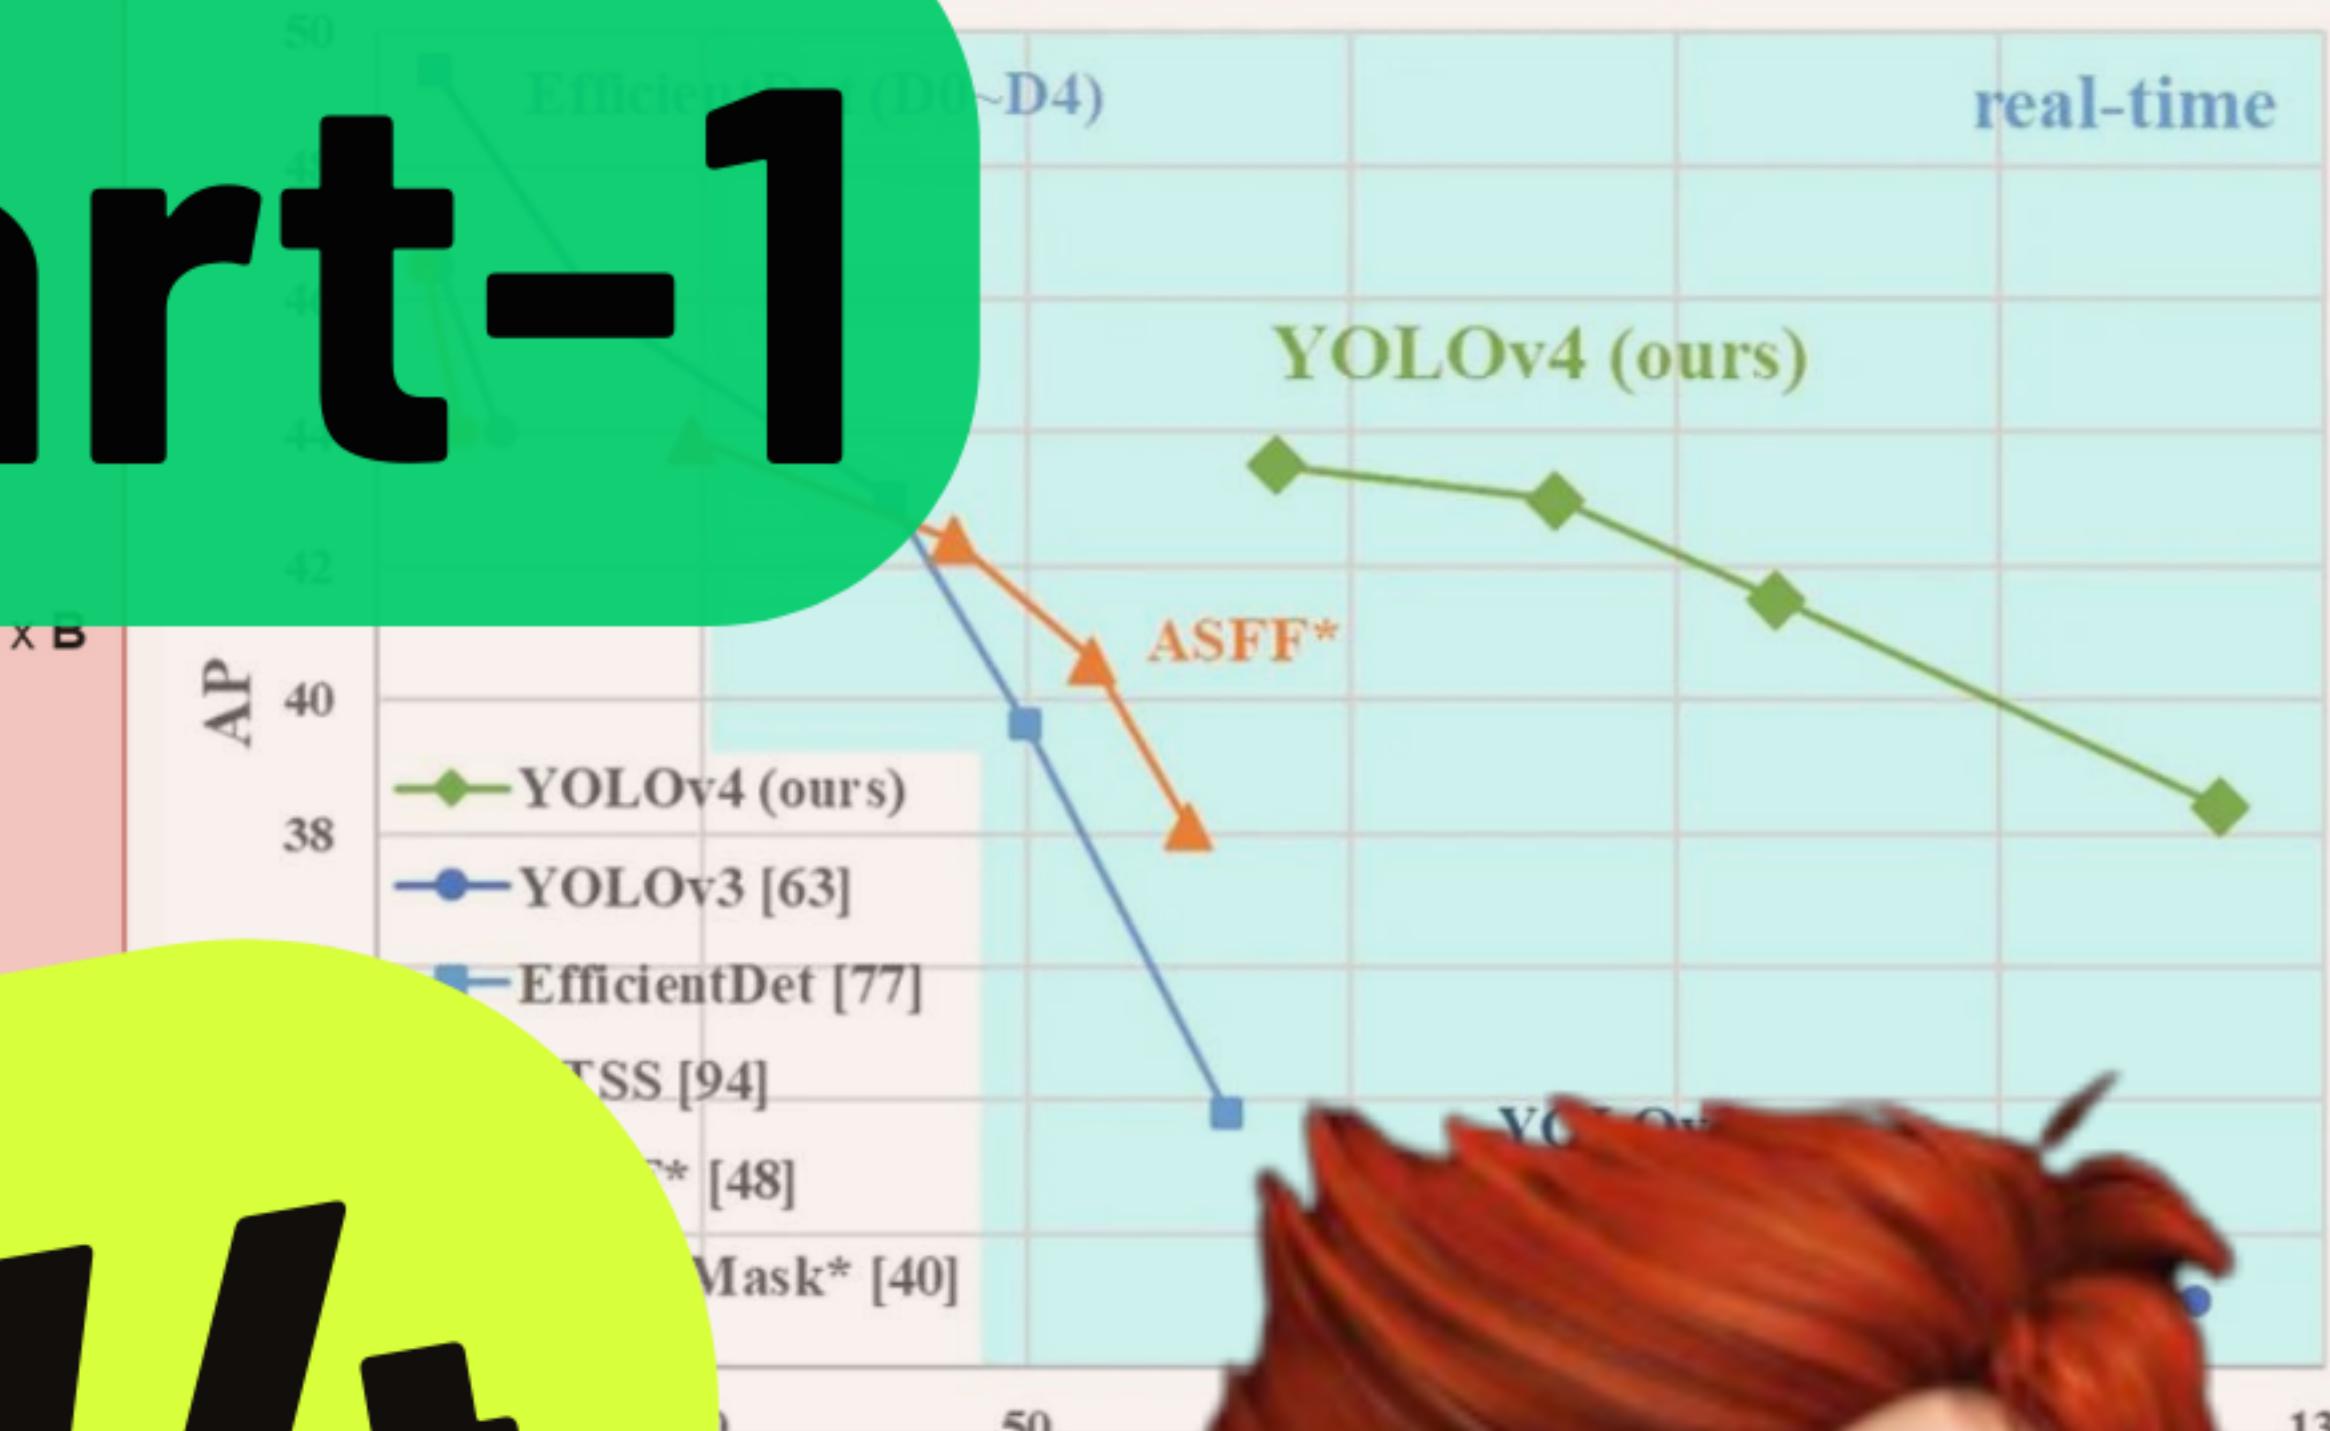
Regularization

- DropOut
- DropPath
- Spatial DropOut
- DropBlock

- IoU
- GIoU
- CIoU
- DIoU
- Distance

MS COCO Object Detection

Part-1



YOLOv3: An Incremental Improvement

Joseph Redmon Ali Farhadi
University of Washington

Abstract

We present some updates to YOLO! We made a bunch of little design changes to make it better. We also trained this new network that's pretty swell. It's a little bigger than last time but more accurate. It's still fast though, don't worry. At 320×320 YOLOv3 runs in 22 ms at 28.2 mAP, as accurate as SSD but three times faster. When we look at the old .5 IOU mAP detection metric YOLOv3 is quite good. It achieves 57.9 AP₅₀ in 51 ms on a Titan X, compared to 57.5 AP₅₀ in 198 ms by RetinaNet, similar performance but 3.8× faster. As always, all the code is online at <https://pjreddie.com/yolo/>.

1. Introduction

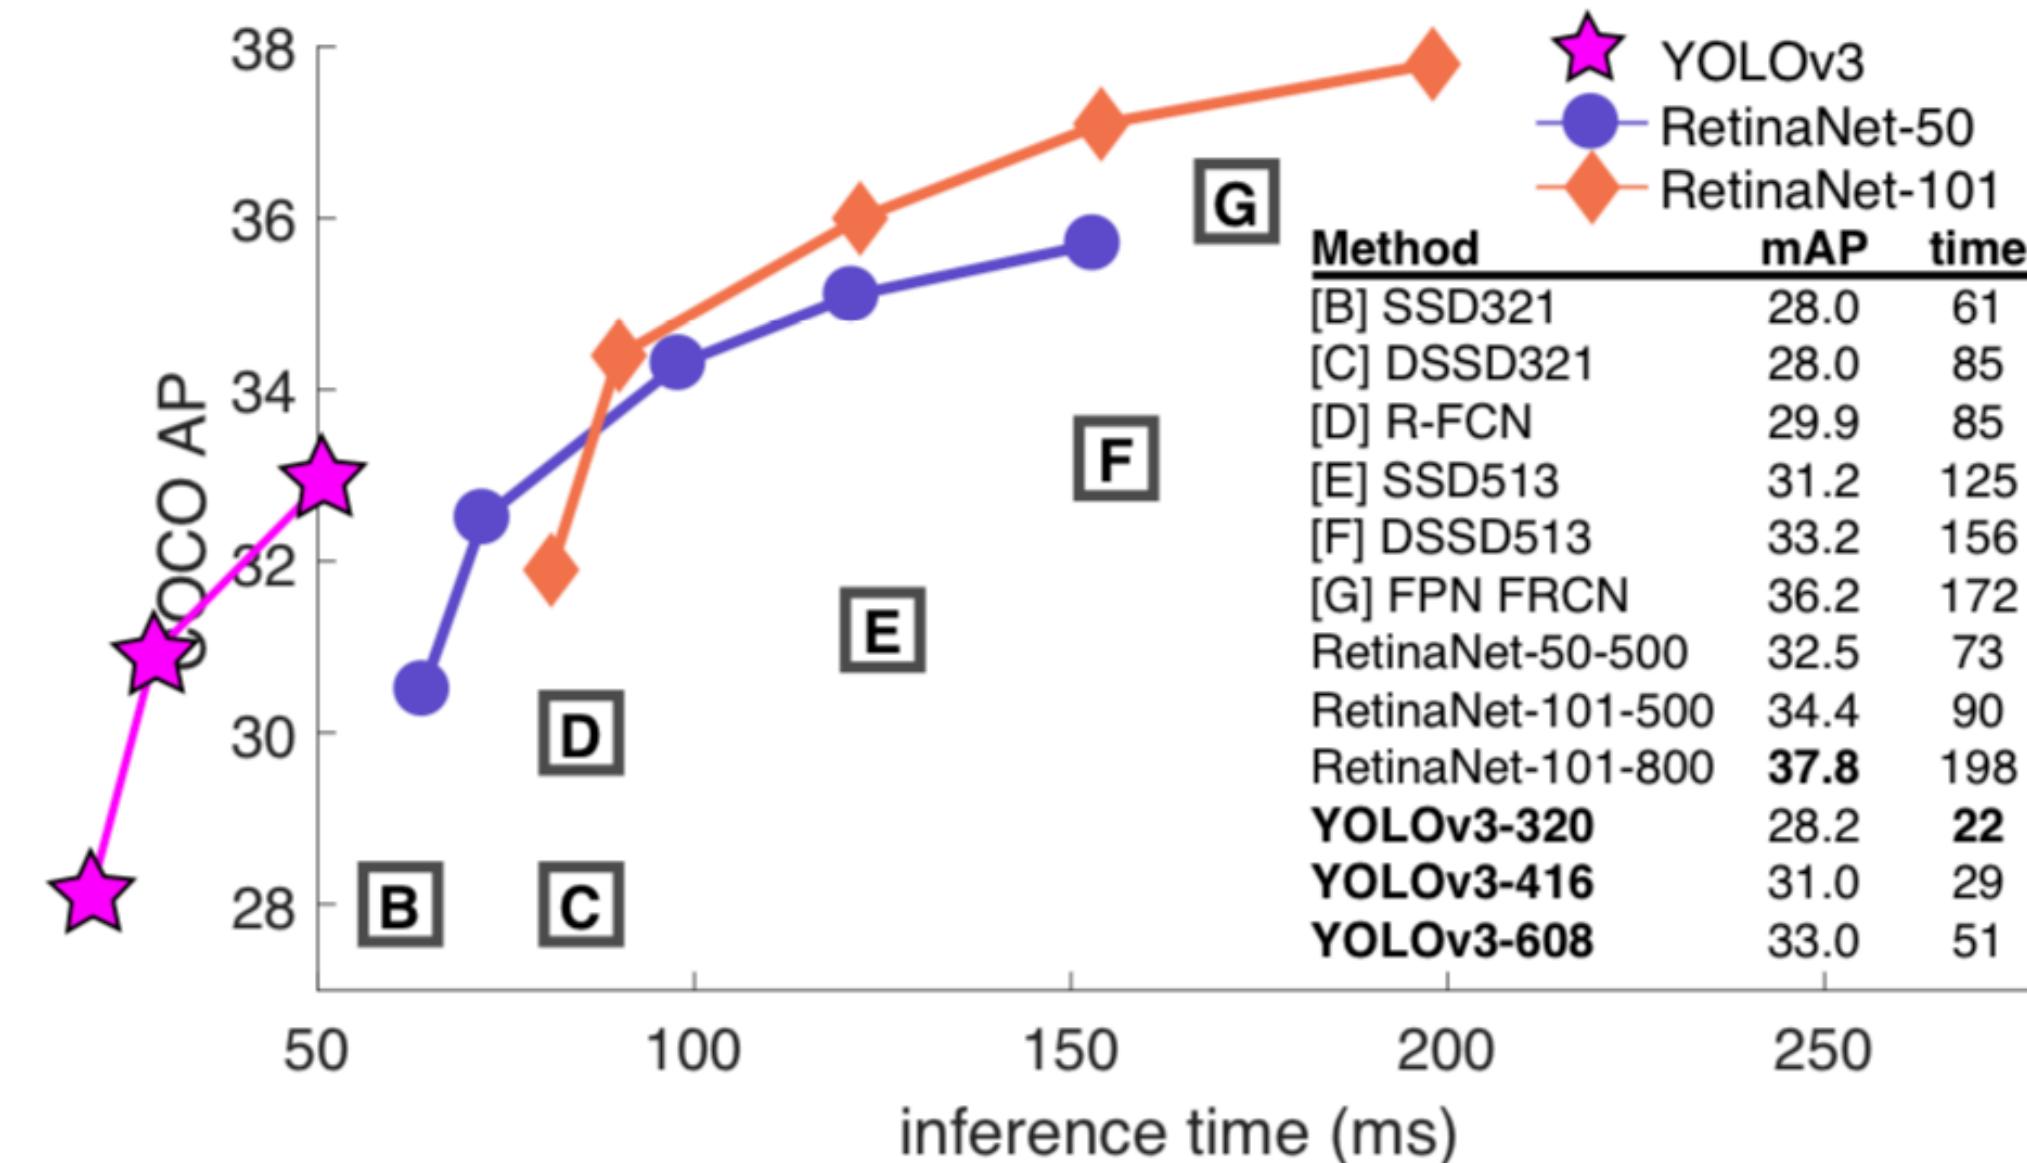


Figure 1. We adapt this figure from the Focal Loss paper [9]. YOLOv3 runs significantly faster than other detection methods with comparable performance. Times from either an M40 or Titan X, they are basically the same GPU.



Joe Redmon
@pjreddie

I stopped doing CV research because I saw the impact my work was having. I loved the work but the military applications and privacy concerns eventually became impossible to ignore.[twitter.com/RogerGrosse/st...](https://twitter.com/RogerGrosse/status/1034311100000000000)



Alexey

AlexeyAB

[Follow](#)

Alexey Bochkovskiy (Aleksei Bochkovskii)

4.4k followers · 4 following

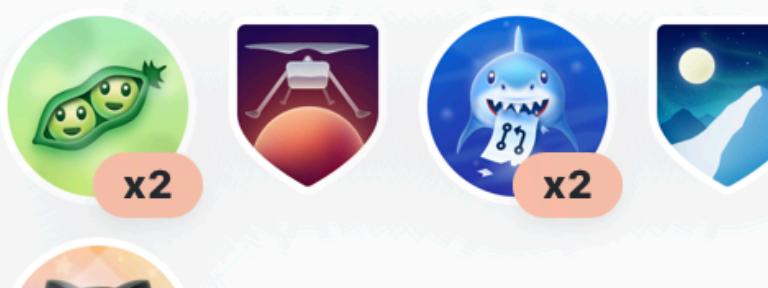
Ex Intel

Germany, Munich

<https://scholar.google.com/citations?user=ljmswJ0AAAAJ>

@alexeyab84

Achievements



Overview Repositories 123 Projects Packages 52 Stars 699

darknet Public

Forked from pjreddie/darknet

YOLOv4 / Scaled-YOLOv4 / YOLO - Neural Networks for Object Detection (Windows and Linux version of Darknet)

C 20.3k 7.8k

WongKinYiu/yolov7 Public

Implementation of paper - YOLOv7: Trainable bag-of-freebies sets new state-of-the-art for real-time object detectors

Jupyter Notebook 9.1k 2.8k

isl-org/DPT Public

Dense Prediction Transformers

Python 1.4k 202

object_threadsafe Public

We make any object thread-safe and std::shared_mutex 10 times faster to achieve the speed of lock-free algorithms on >85% reads

C++ 434 114

isl-org/MiDaS Public

Code for robust monocular depth estimation described in "Ranftl et. al., Towards Robust Monocular Depth Estimation: Mixing Datasets for Zero-shot Cross-dataset Transfer, TPAMI 2022"

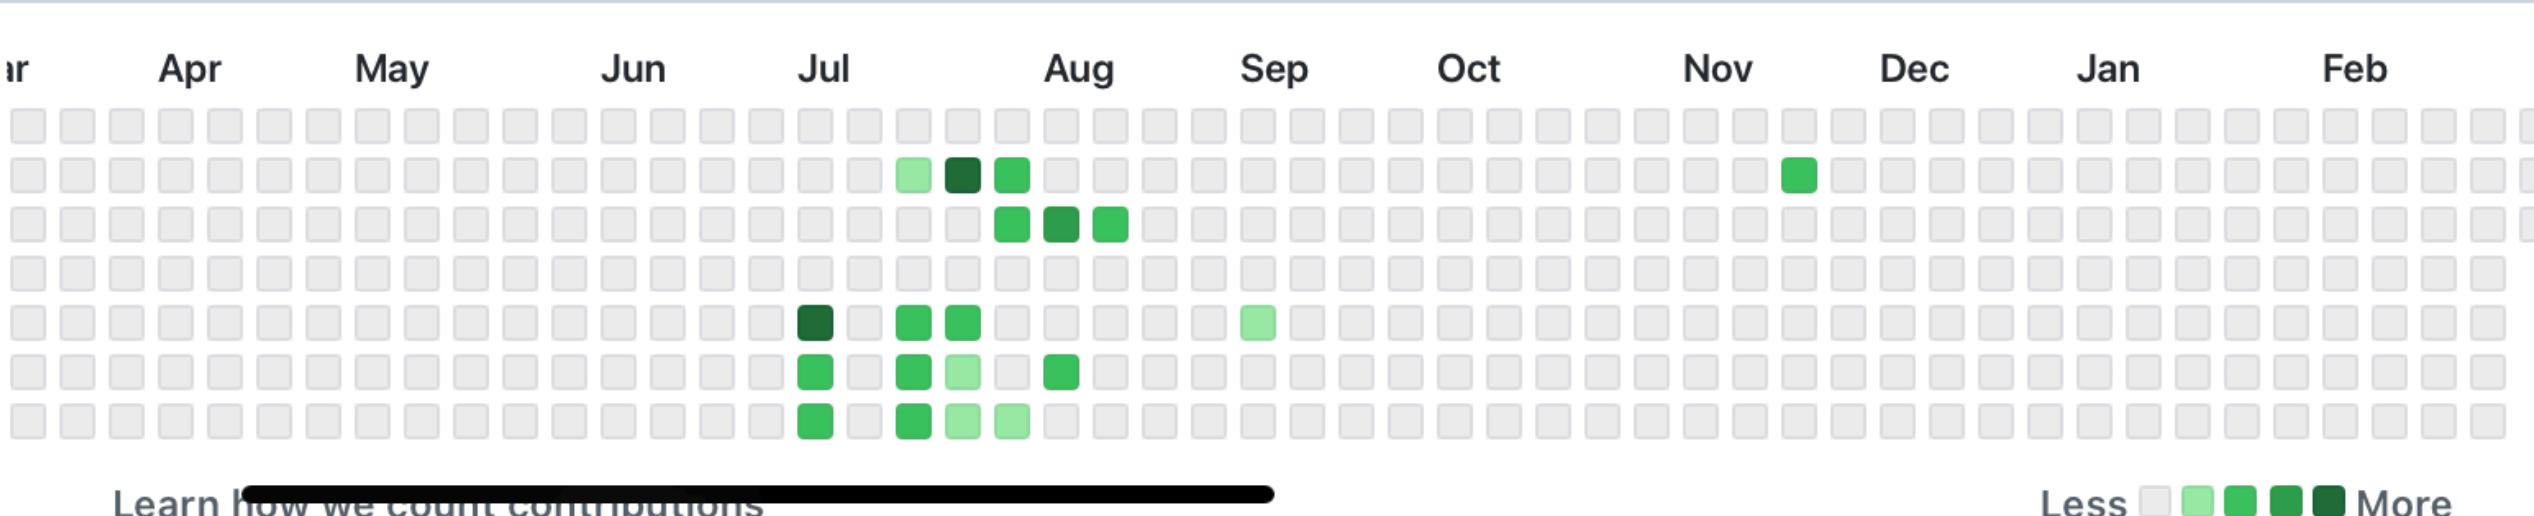
Python 2.6k 408

Yolo_mark Public

GUI for marking bounded boxes of objects in images for training neural network Yolo v3 and v2

C++ 1.7k 670

46 contributions in the last year



YOLOv4: Optimal Speed and Accuracy of Object Detection

Alexey Bochkovskiy*

alexeyab84@gmail.com

Chien-Yao Wang*

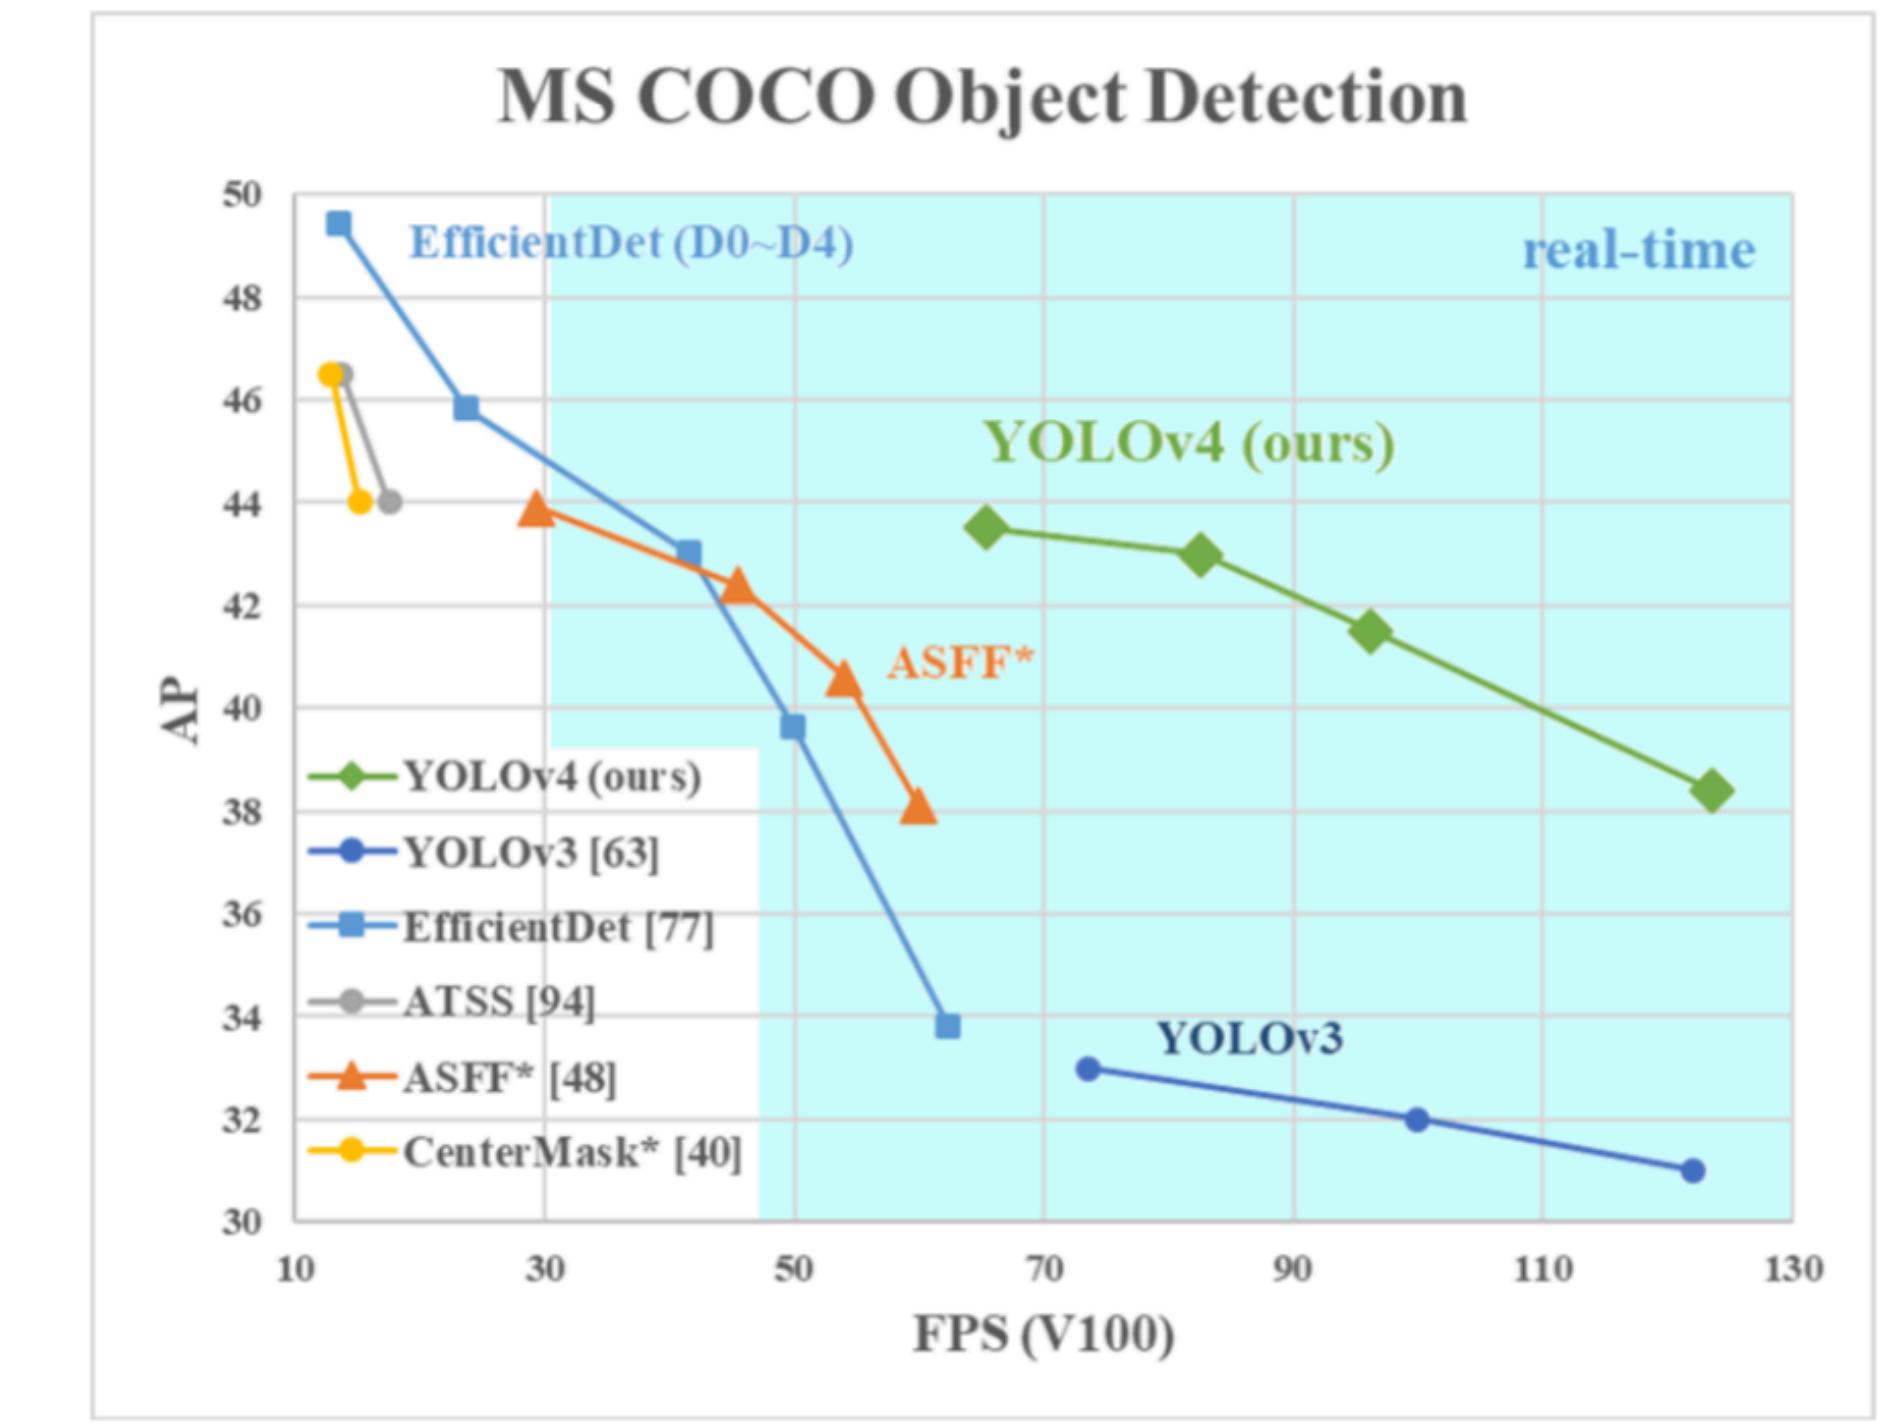
Institute of Information Science
Academia Sinica, Taiwan
kinyiu@iis.sinica.edu.tw

Hong-Yuan Mark Liao

Institute of Information Science
Academia Sinica, Taiwan
liao@iis.sinica.edu.tw

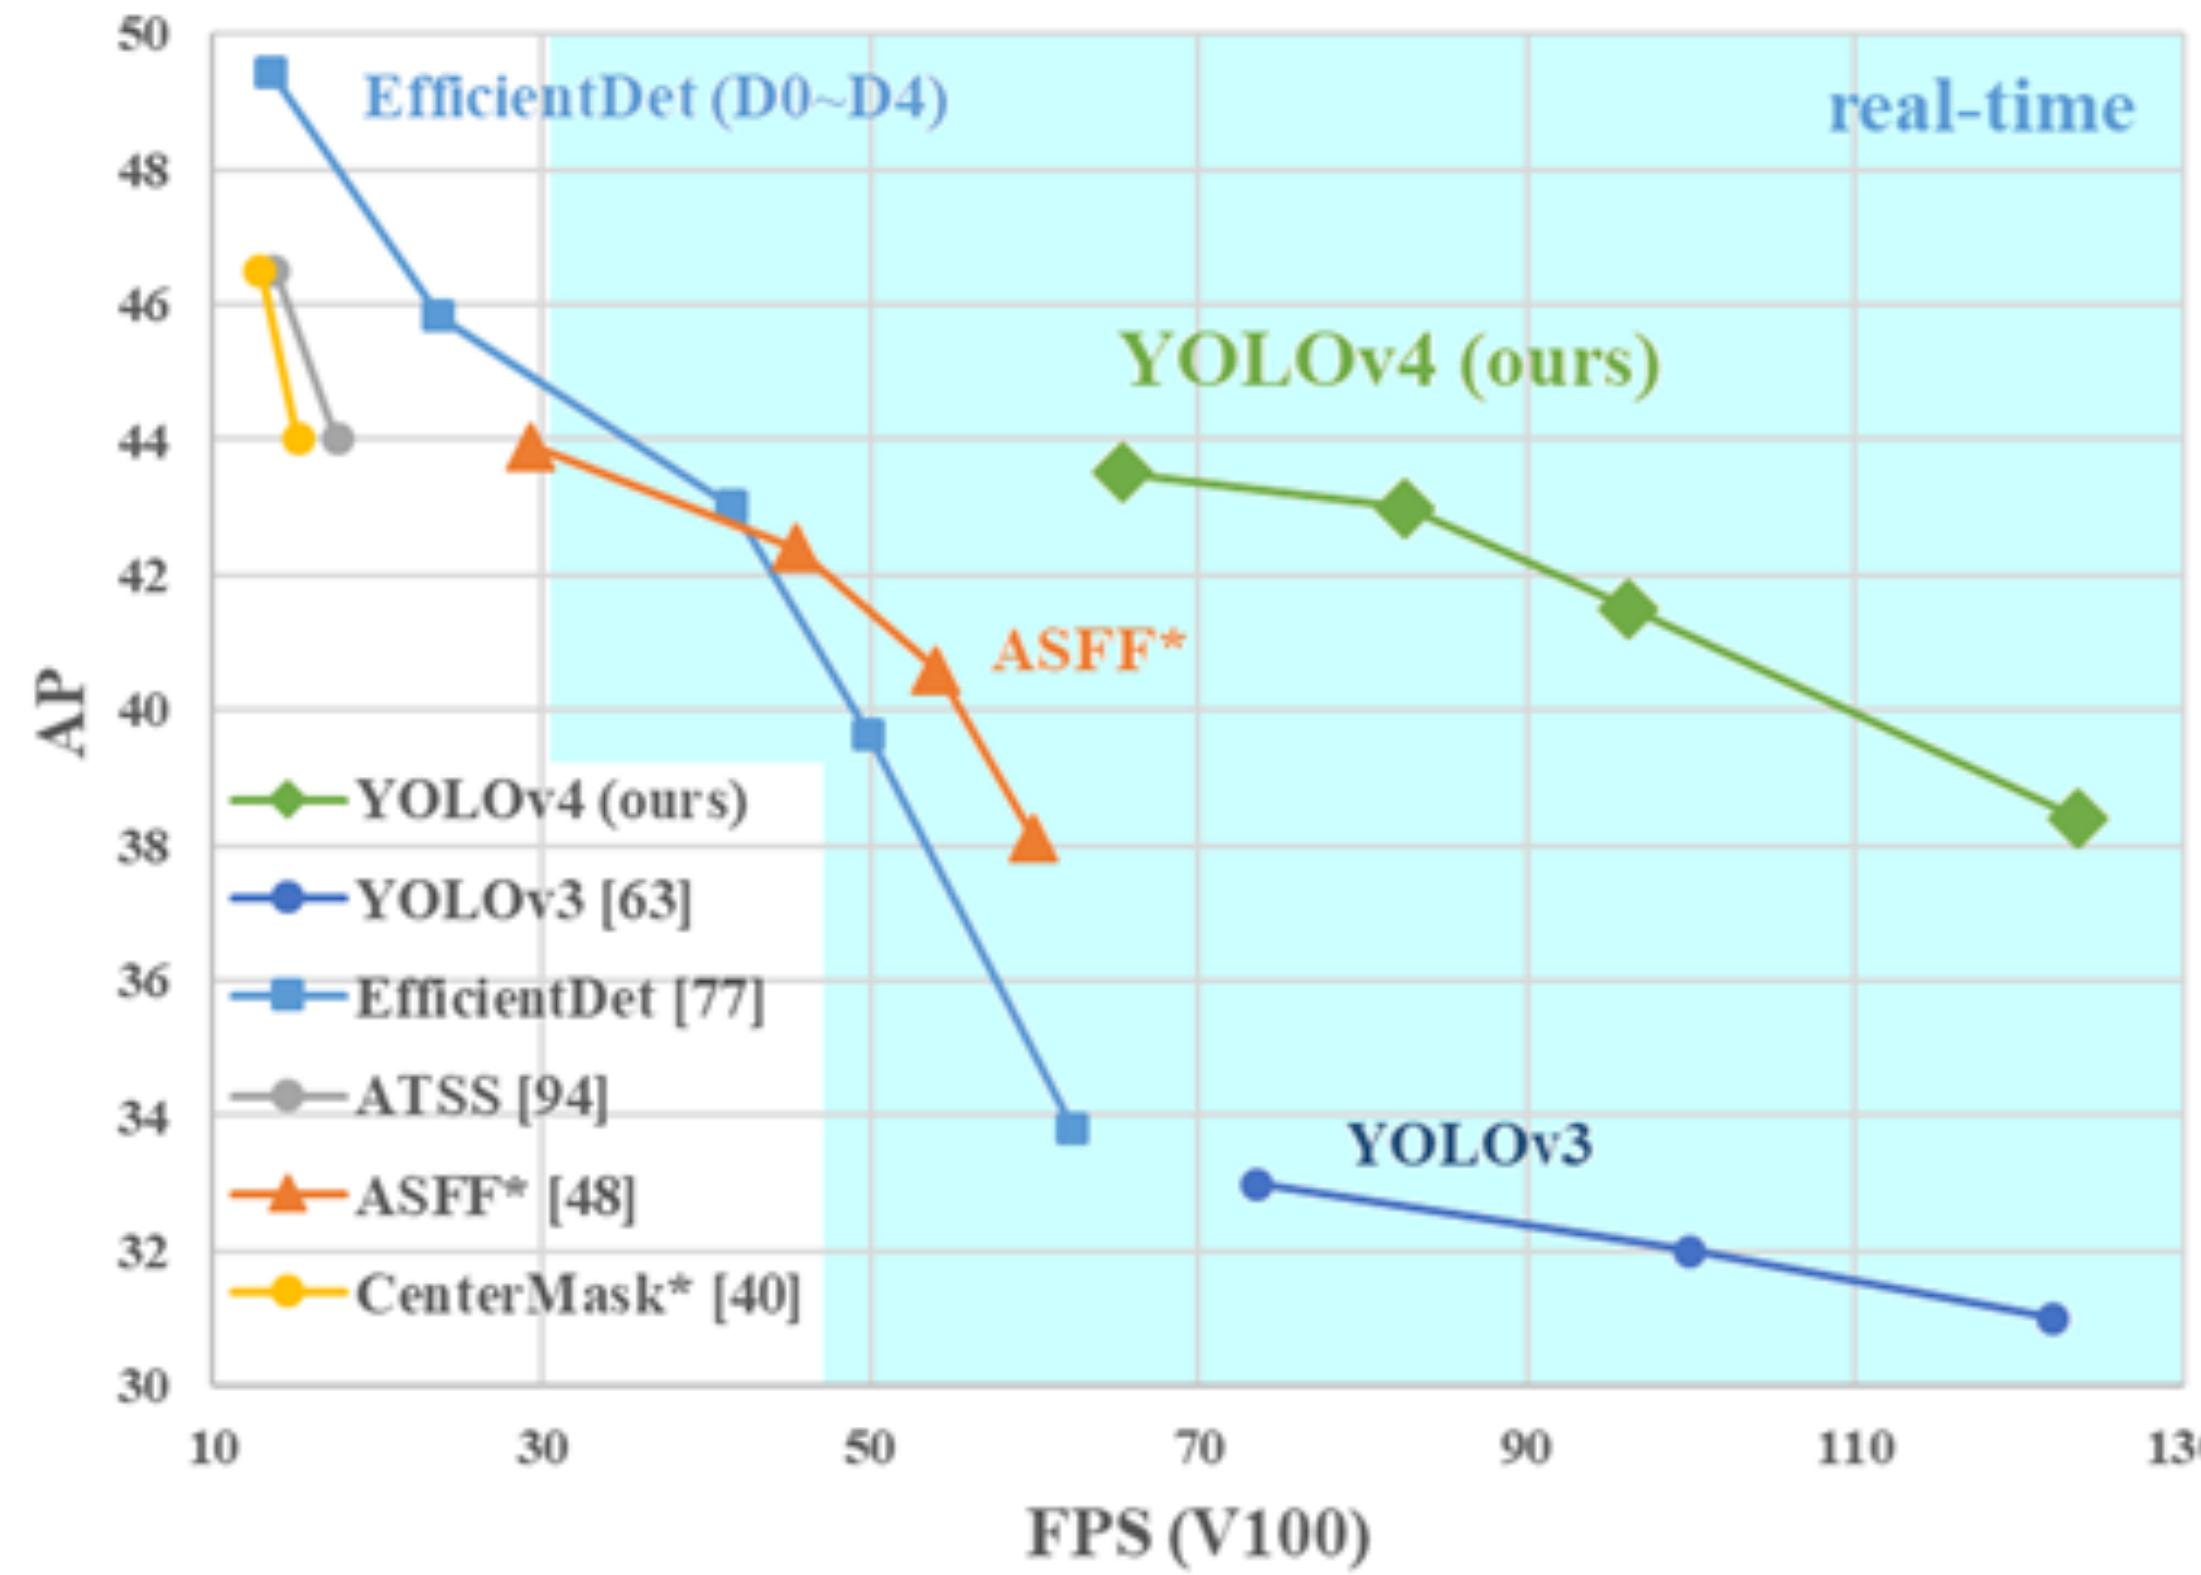
Abstract

There are a huge number of features which are said to improve Convolutional Neural Network (CNN) accuracy. Practical testing of combinations of such features on large datasets, and theoretical justification of the result, is required. Some features operate on certain models exclusively and for certain problems exclusively, or only for small-scale datasets; while some features, such as batch-normalization and residual-connections, are applicable to the majority of models, tasks, and datasets. We assume that such universal features include Weighted-Residual-Connections (WRC), Cross-Stage-Partial-connections (CSP), Cross mini-Batch Normalization (CmBN), Self-adversarial-training (SAT) and Mish-activation. We use new features: WRC, CSP, CmBN, SAT, Mish activation, Mosaic data augmentation,

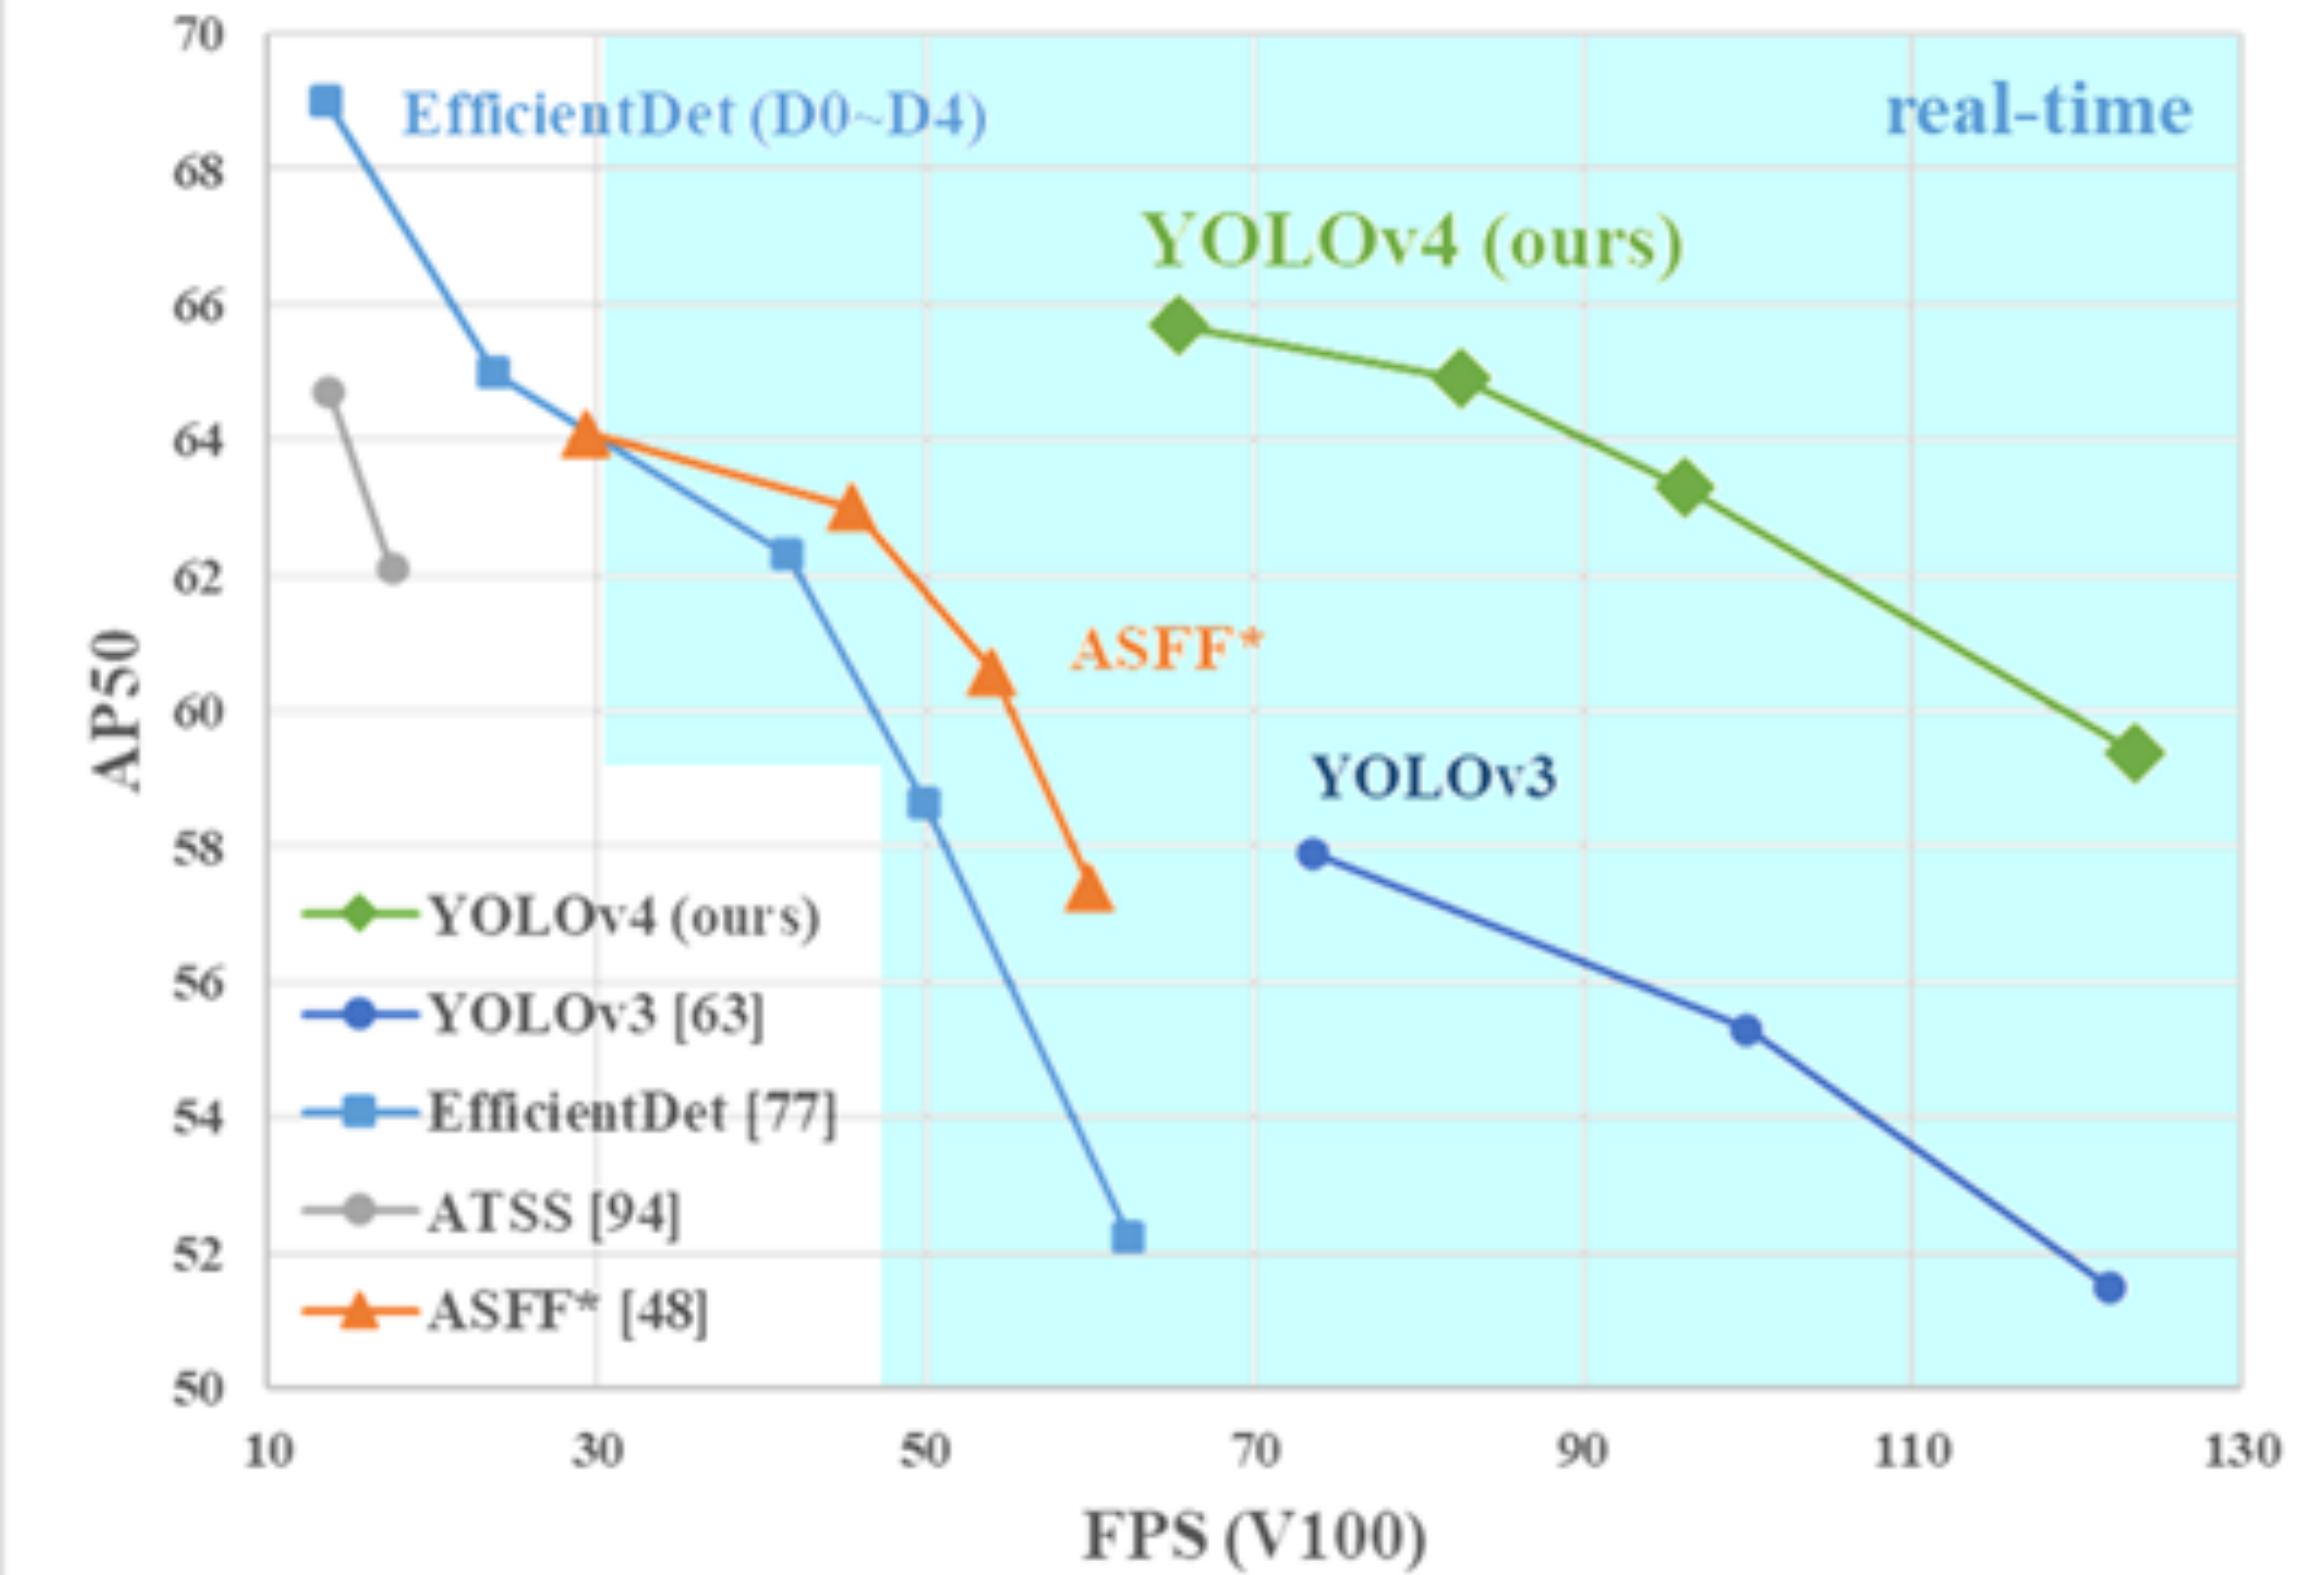


RESULTS

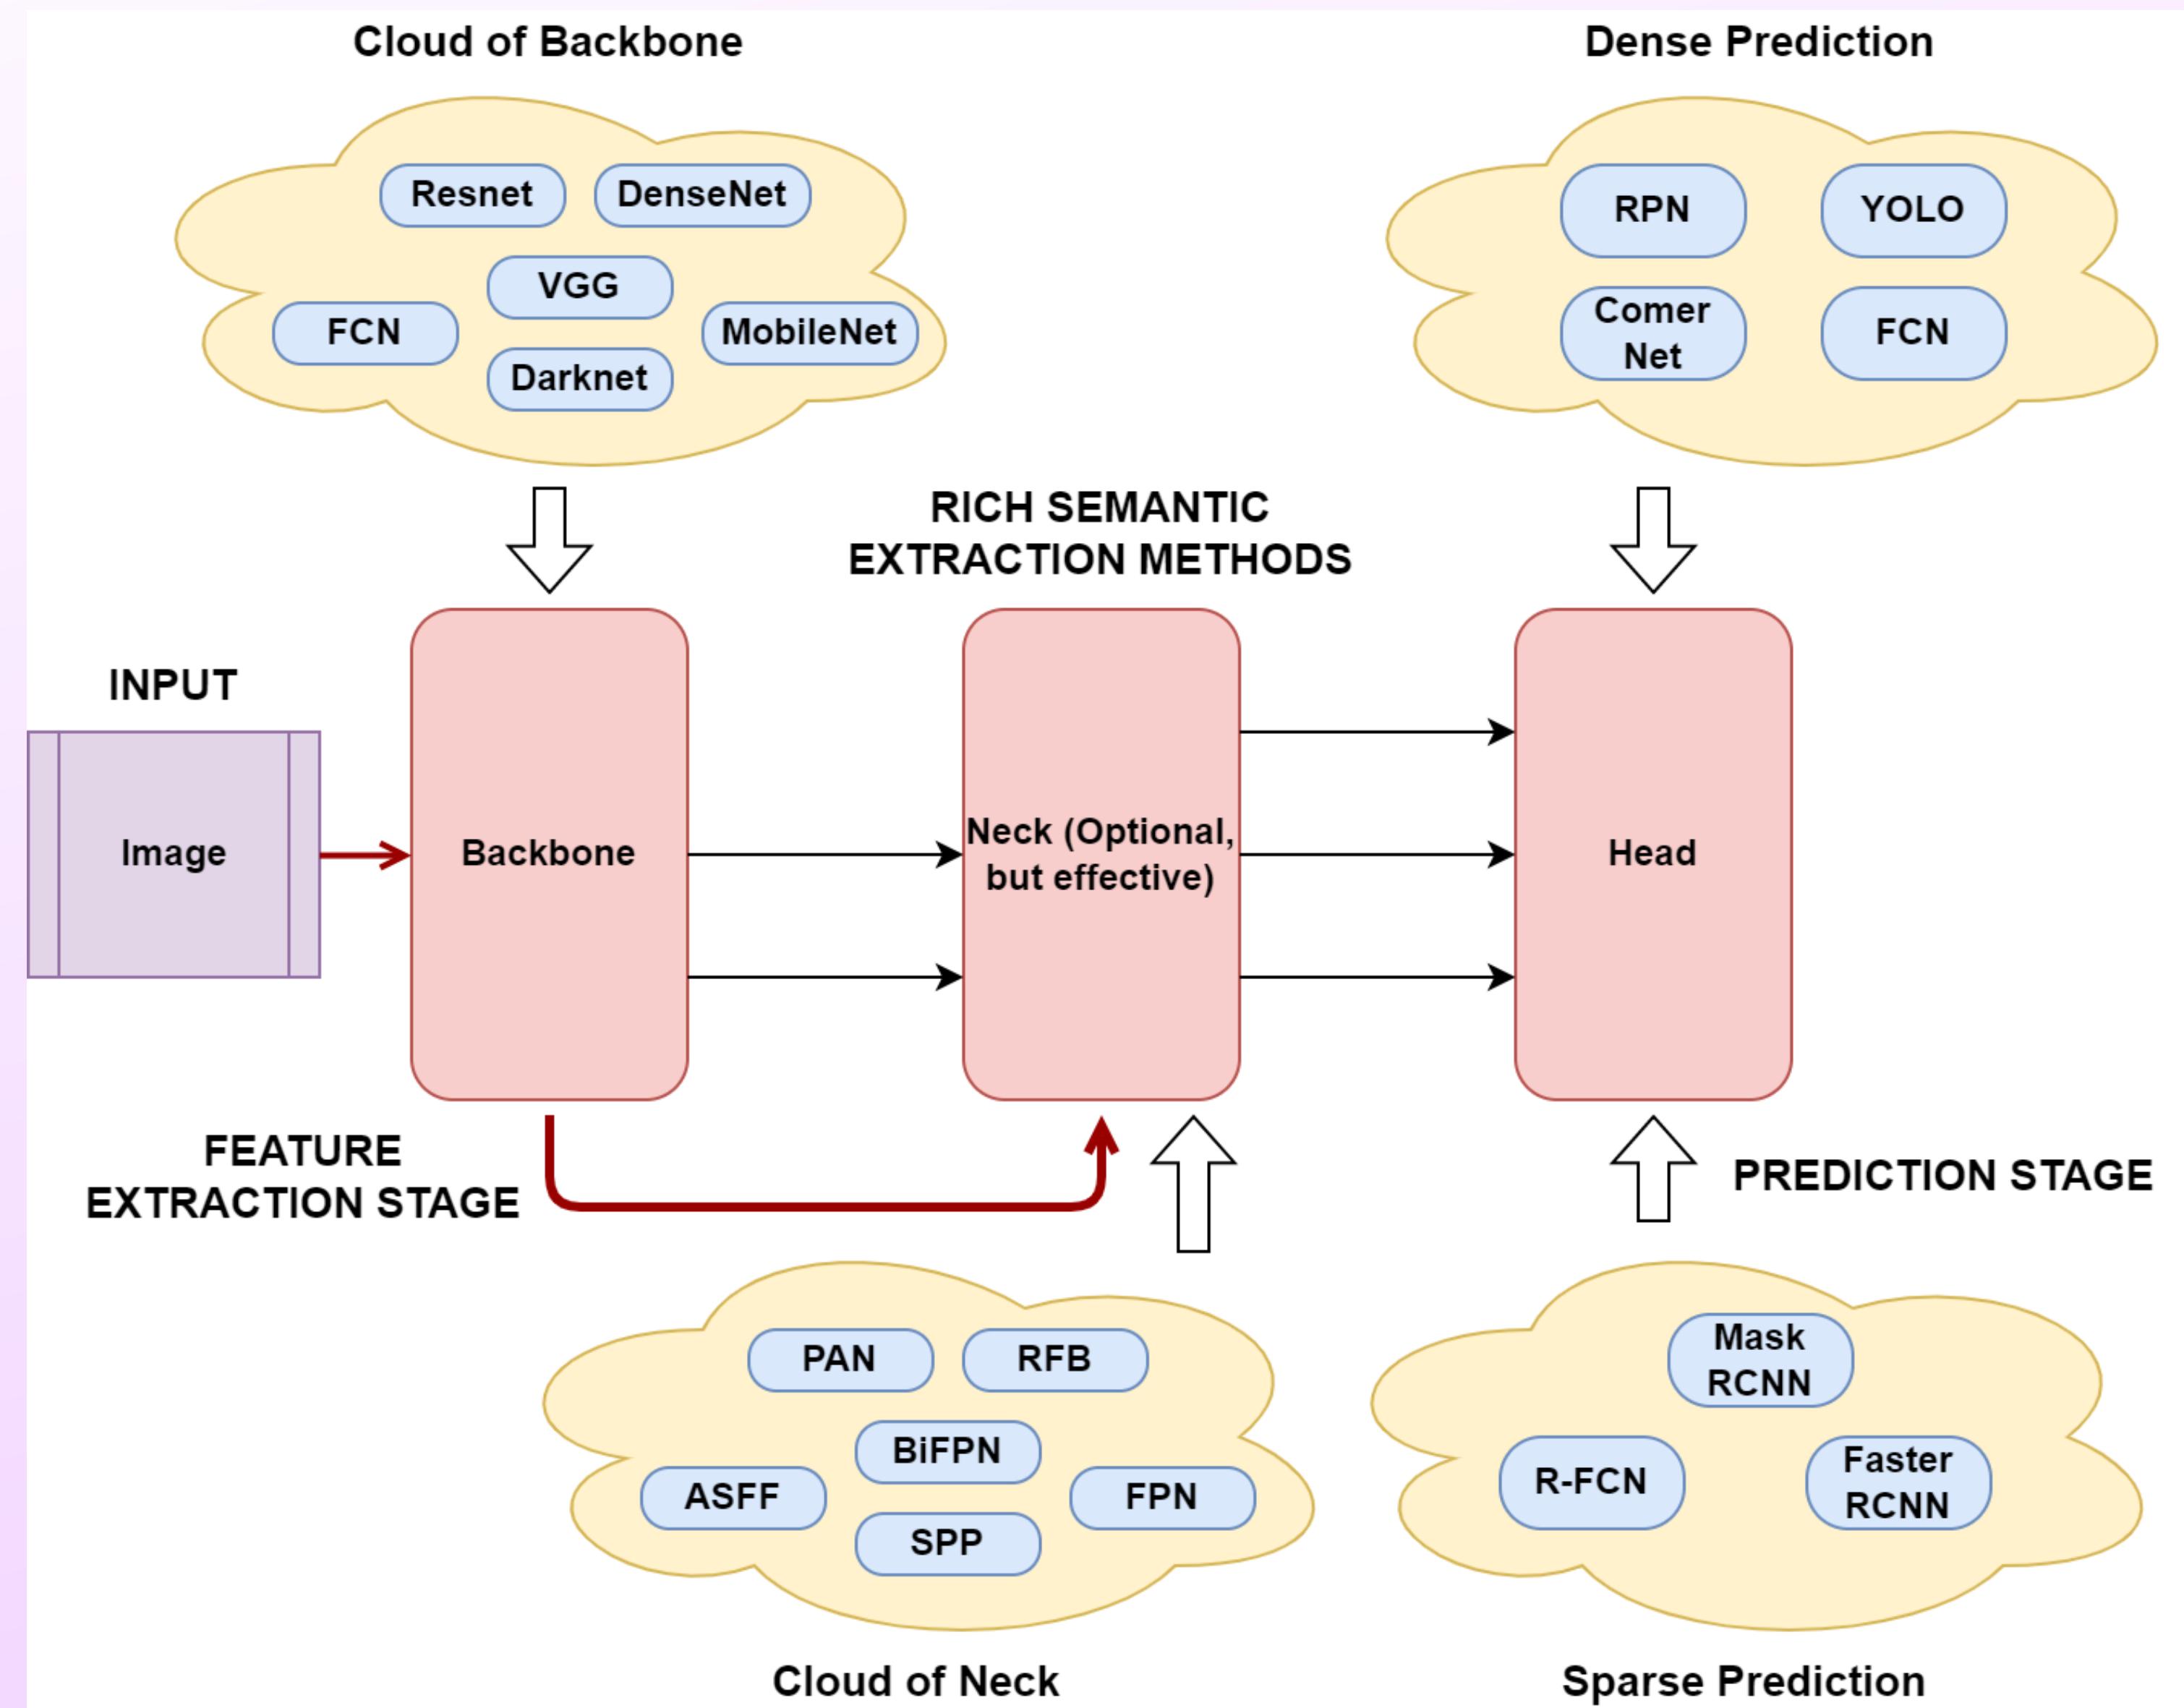
MS COCO Object Detection



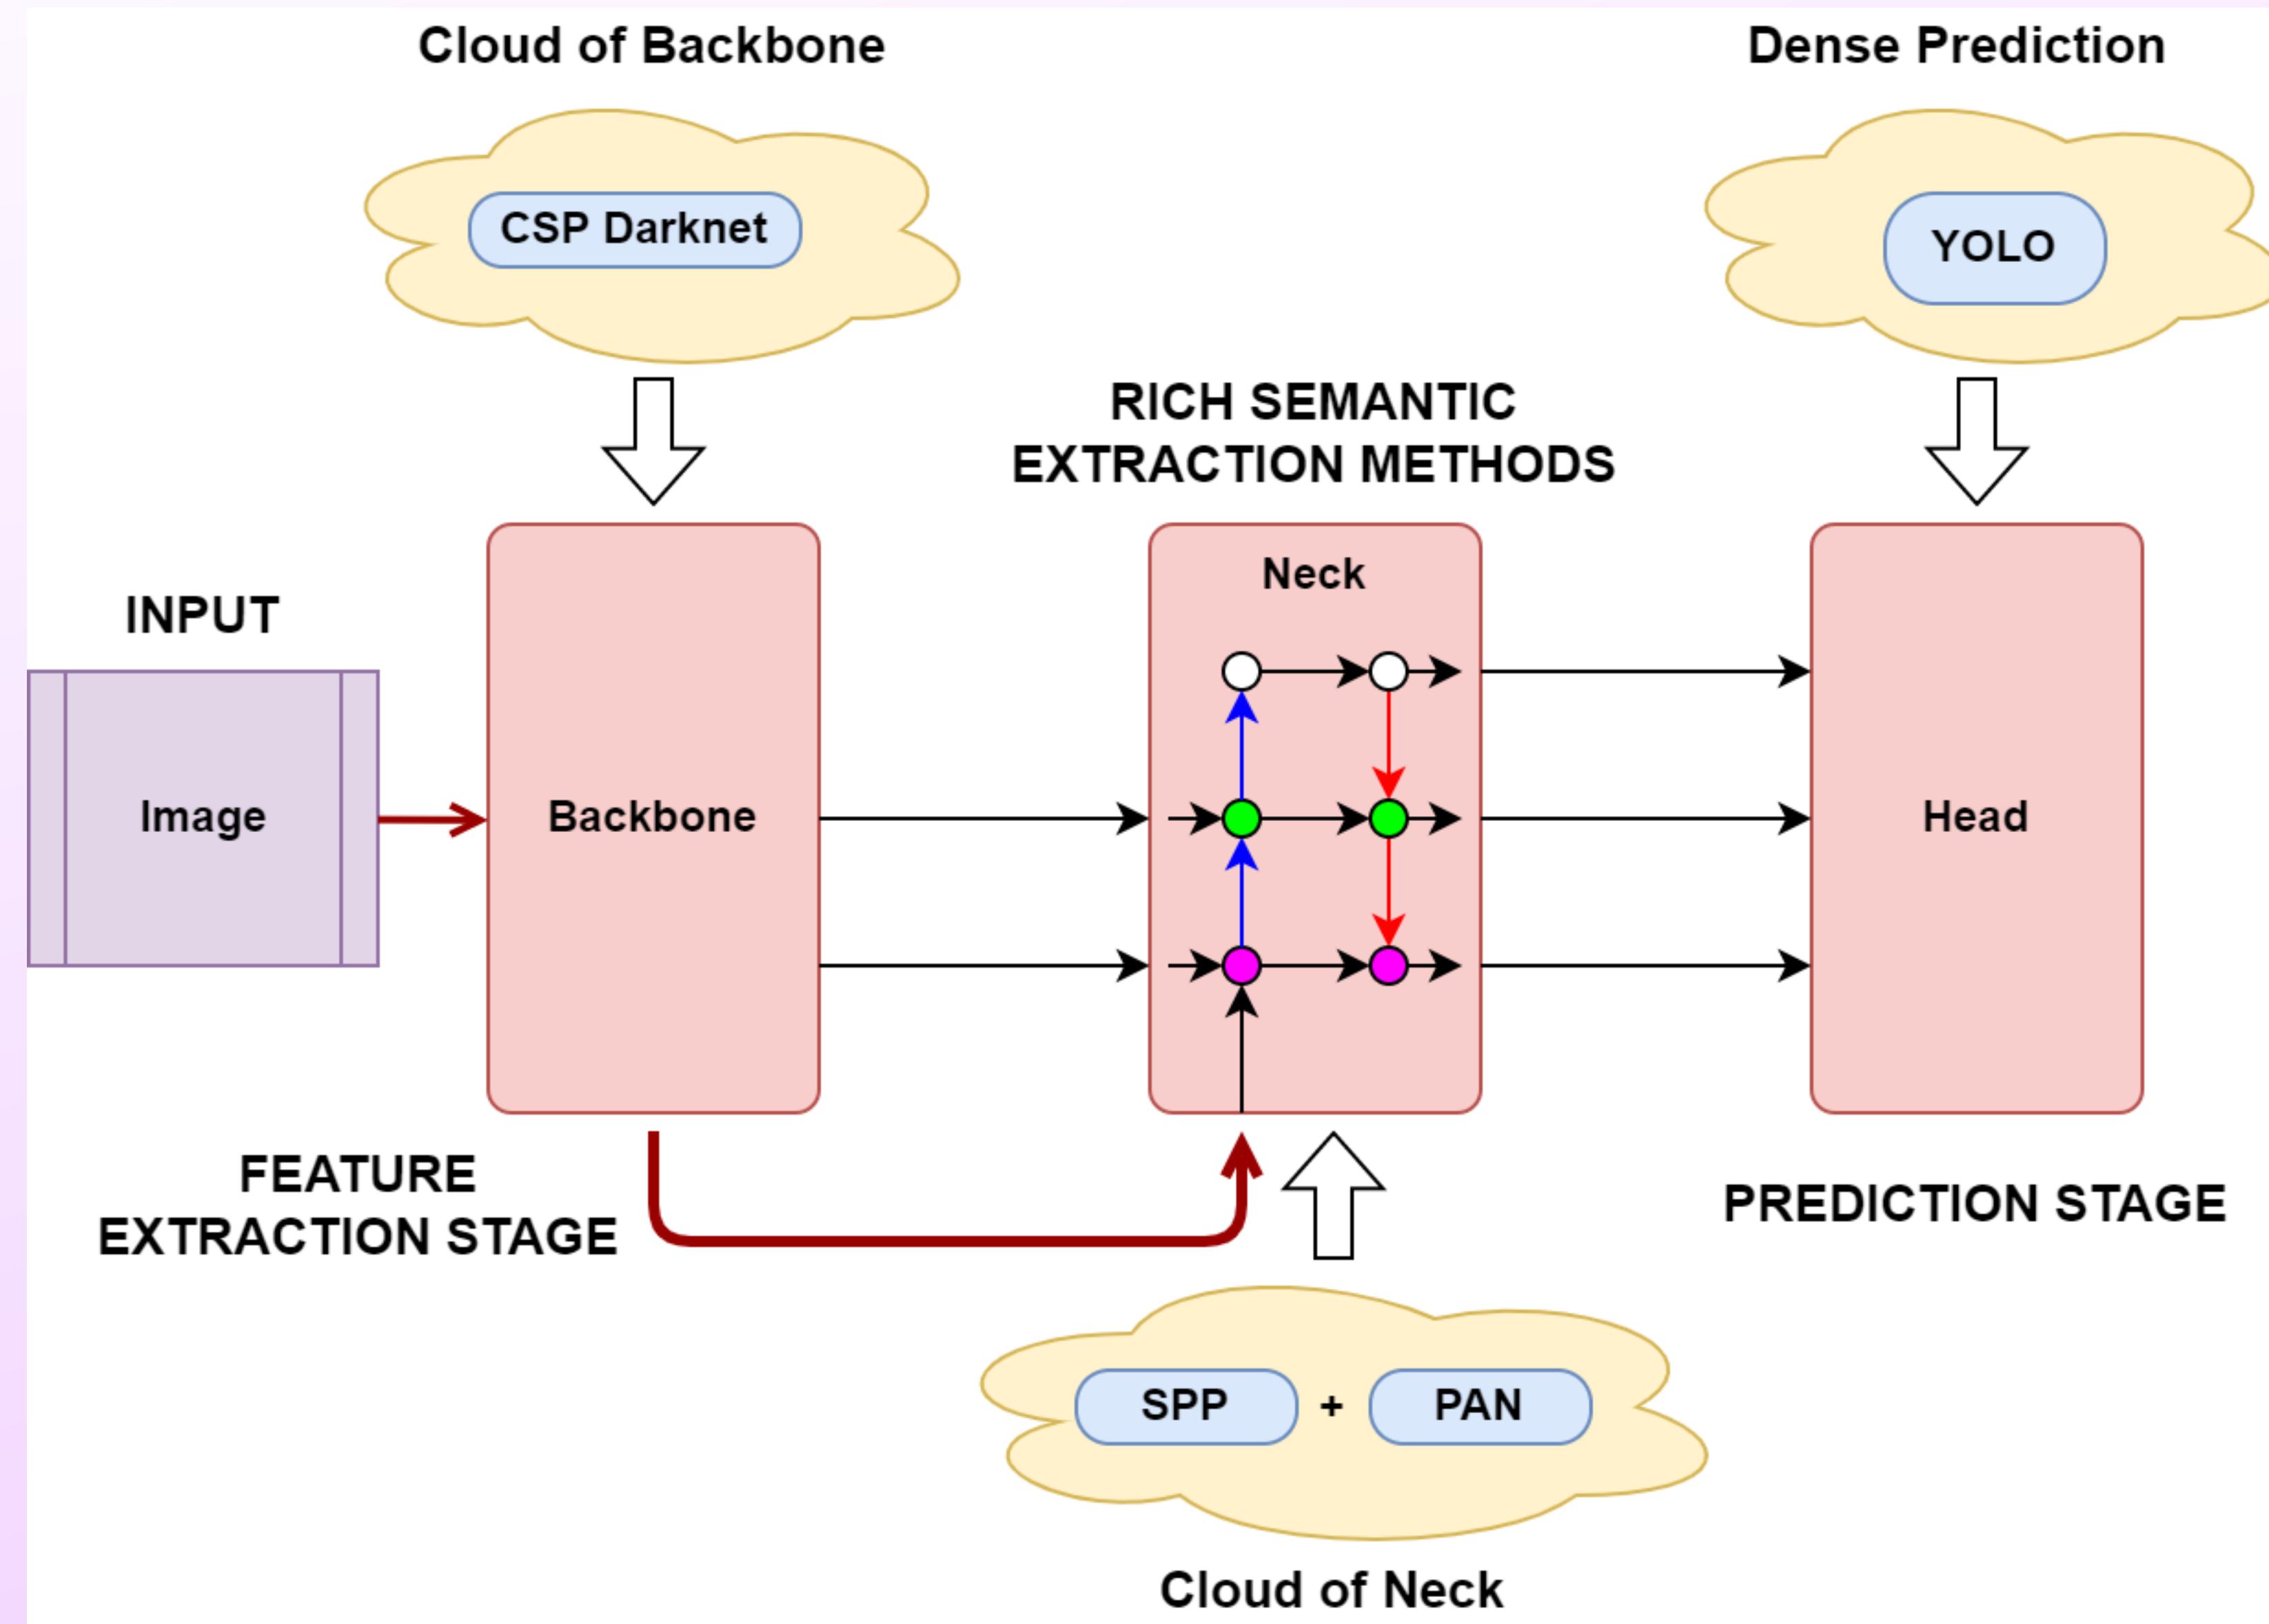
MS COCO Object Detection



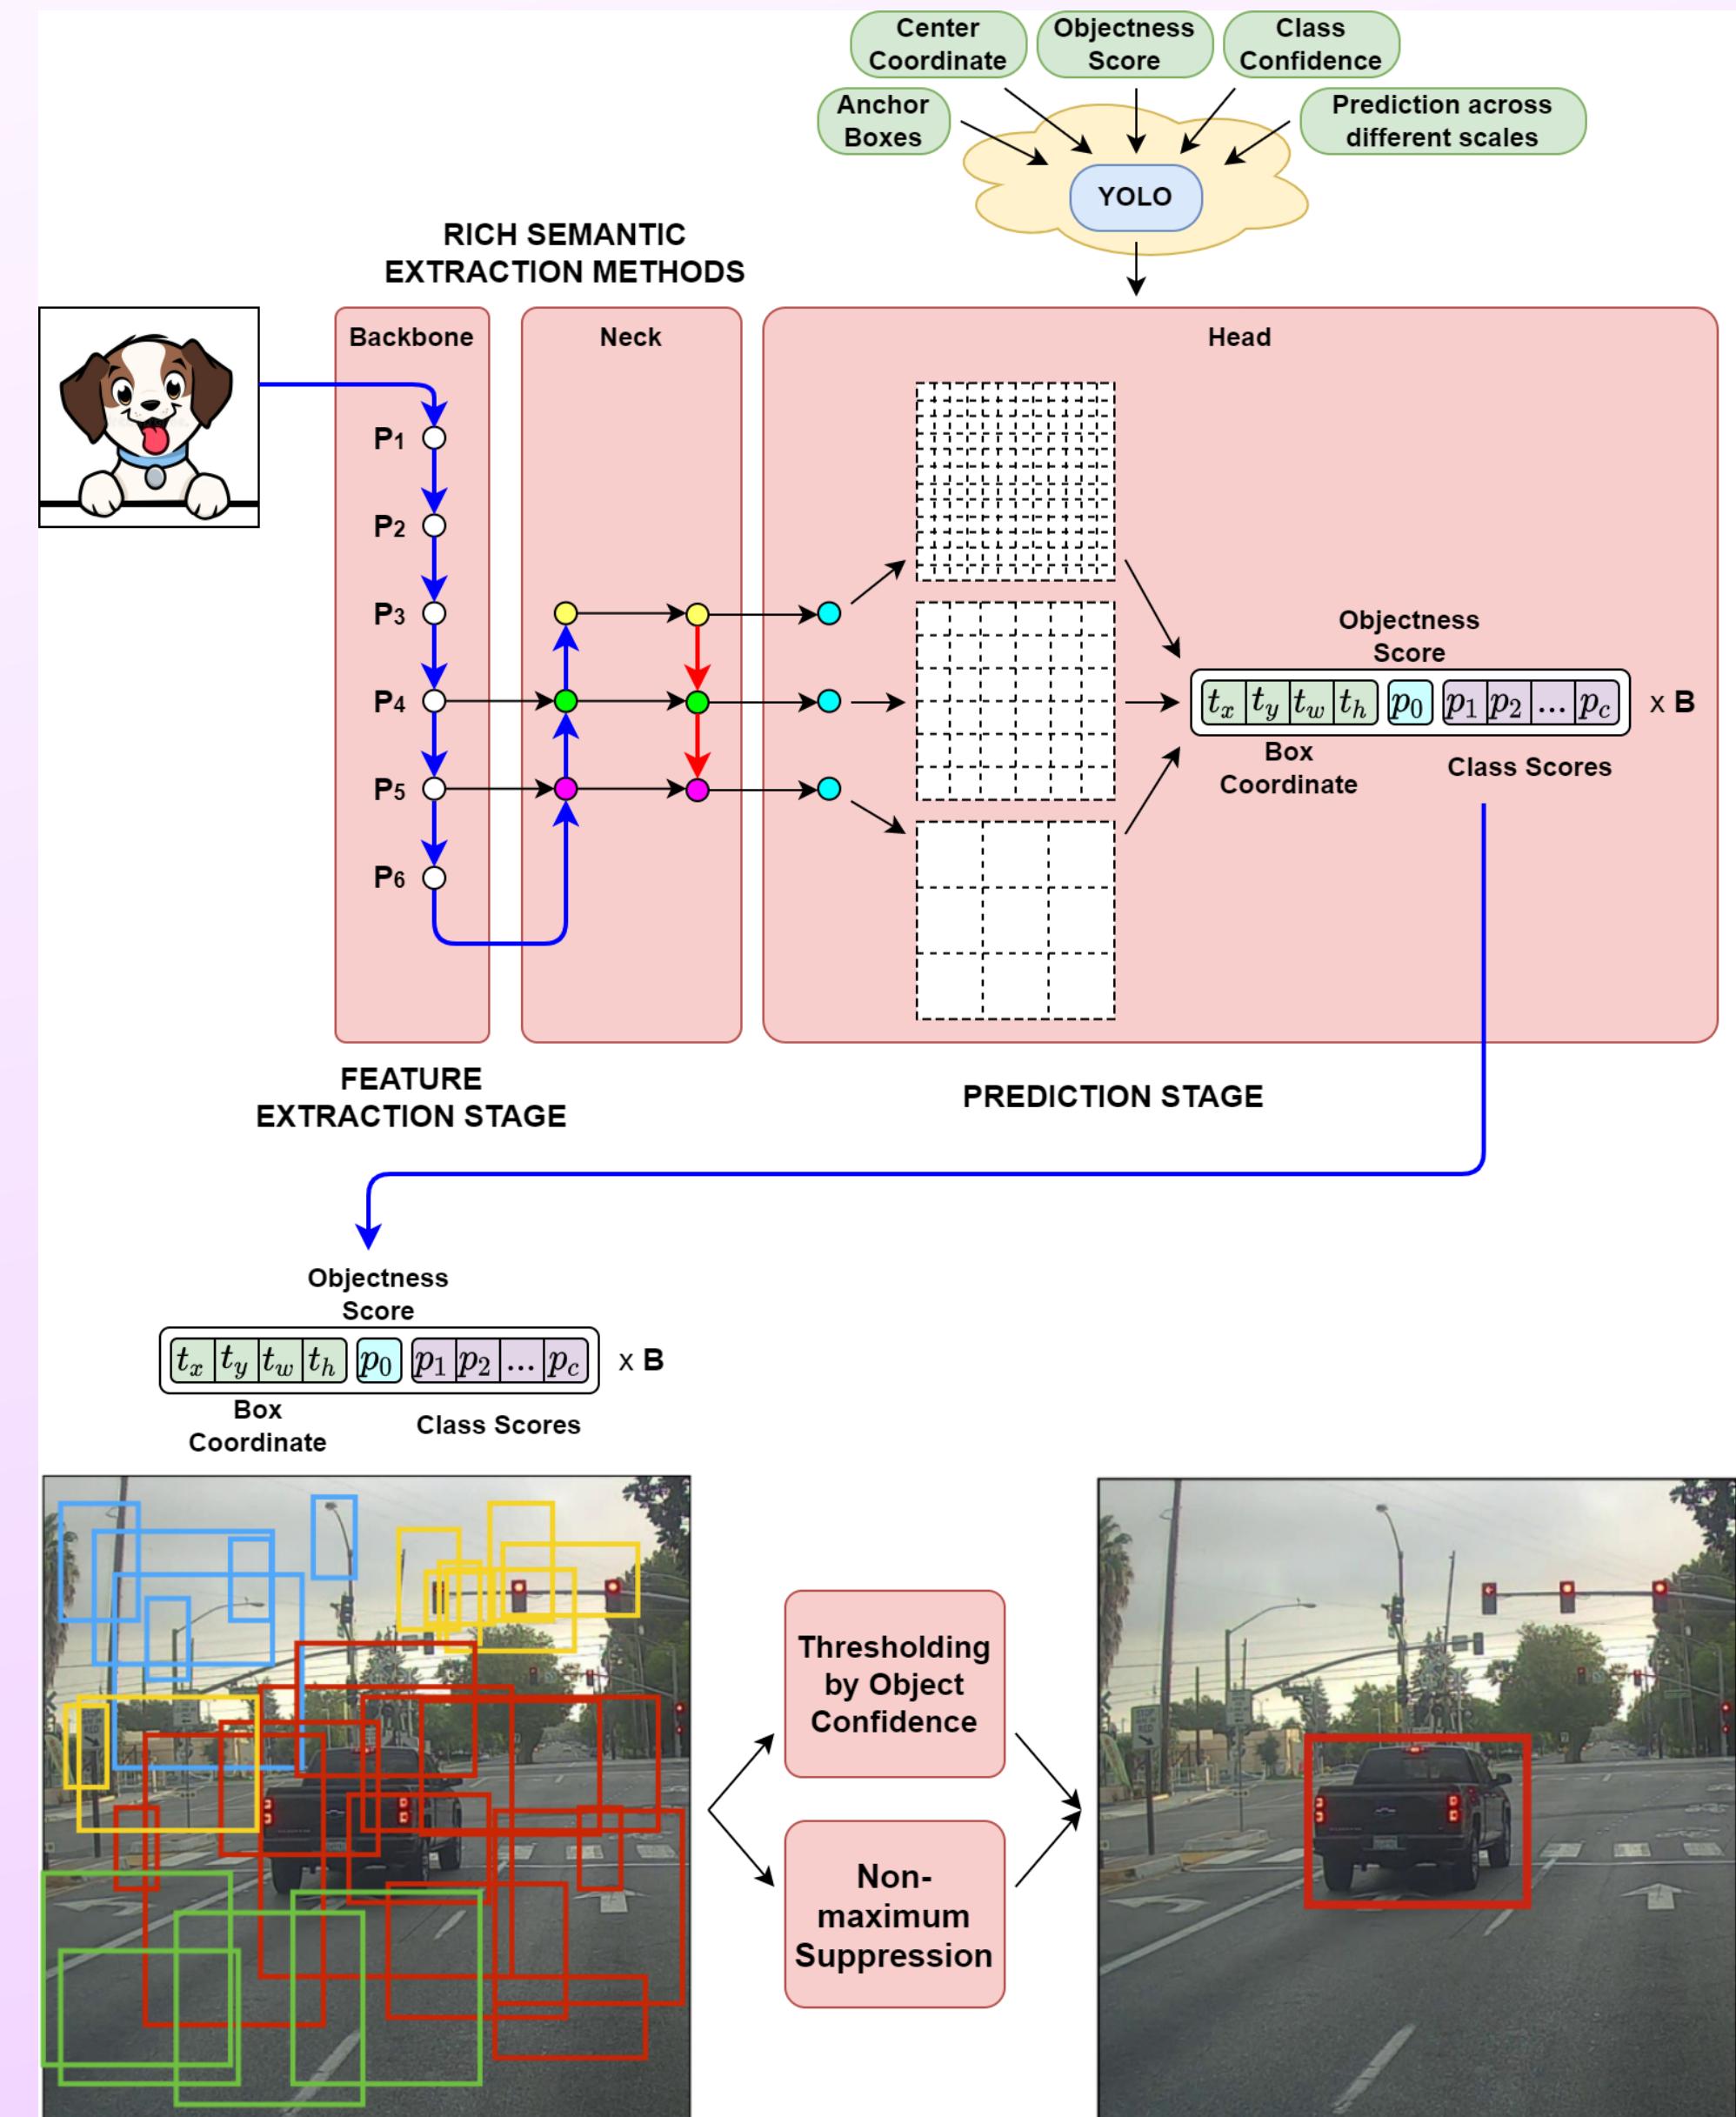
Components



YoloV4



YoloV4



Performance Optimizations

- **Bag of Freebies**
- **Bag of Specials**

Performance Optimizations

- **Bag of Freebies**
 - **Different opt methods during training and inference**
 - **Doesn't increase the inference time**
 - **Increases the training time**
- **Bag of Specials**
 - **Different modules that drastically improve accuracy**
 - **Increase the inference time by a small amount**

BoF & BoS

	Backbone	Detector
Bag of Freebies (BoF)	<ul style="list-style-type: none">• CutMix• Mosaic data augmentation• DropBlock• Class label smoothing	<ul style="list-style-type: none">• CloU-loss• Cross mini-Batch Normalization• DropBlock• Mosaic data augmentation• Self-Adversarial Training• Multiple anchors for a single ground truth• Cosine annealing scheduler• Optimal hyperparameters• Random training shapes
Bag of Specials (BoS)	<ul style="list-style-type: none">• Mish activation• Cross-stage partial connections (CSP)• Multi-input weighted residual connections (MiWRC)	<ul style="list-style-type: none">• Mish activation• SPP-block• SAM-block• PAN path-aggregation block• DIoU-NMS

Techniques Evaluated

- **Input:** Image, Patches, Image Pyramid
- **Backbones:** VGG16 [68], ResNet-50 [26], SpineNet [12], EfficientNet-B0/B7 [75], CSPResNeXt50 [81], CSPDarknet53 [81]
- **Neck:**
 - **Additional blocks:** SPP [25], ASPP [5], RFB [47], SAM [85]
 - **Path-aggregation blocks:** FPN [44], PAN [49], NAS-FPN [17], Fully-connected FPN, BiFPN [77], ASFF [48], SFAM [98]
- **Heads:**
 - **Dense Prediction (one-stage):**
 - RPN [64], SSD [50], YOLO [61], RetinaNet [45] (anchor based)
 - CornerNet [37], CenterNet [13], MatrixNet [60], FCOS [78] (anchor free)
 - **Sparse Prediction (two-stage):**
 - Faster R-CNN [64], R-FCN [9], Mask R-CNN [23] (anchor based)
 - RepPoints [87] (anchor free)
- **Activations:** ReLU, leaky-ReLU, parametric-ReLU, ReLU6, SELU, Swish, or Mish
- **Bounding box regression loss:** MSE, IoU, GIoU, CIoU, DIoU
- **Data augmentation:** CutOut, MixUp, CutMix
- **Regularization method:** DropOut, DropPath [36], Spatial DropOut [79], or DropBlock
- **Normalization of the network activations by their mean and variance:** Batch Normalization (BN) [32], Cross-GPU Batch Normalization (CGBN or SyncBN) [93], Filter Response Normalization (FRN) [70], or Cross-Iteration Batch Normalization (CBN) [89]
- **Skip-connections:** Residual connections, Weighted residual connections, Multi-input weighted residual connections, or Cross stage partial connections (CSP)

Thank you!